

hw1

September 13, 2024

0.0.1 Woodrow Reese

HW1

```
[1]: import numpy as np
import pandas as pd
import math
from collections import Counter
```

1 Question 4

1.1 (a)

- Fraud detection - classification. Predictors: Purchase history, recipient's info, transaction suze
- Risk management - classification. Predictors: credit score, debt, cash flow
- Image classification (medical) ## (b)
- My linear regression model in Question 9. Predicting MPG given car weight.
- Population change over time. Predictors: Rent, income, transportation, crime.
- Home valuation. Predictors: SqFt, land, garage, bathrooms, bedrooms. The response is the estimated value of the home given X ## (c)
- Netflix recommendation system. Recommendations are based on what other people who have watched/liked the same show as you.
- Ad targeting. An audience is targeted based on interests, age, search activity etc. Different ads can catch the eye better given a different audiences.
- Player efficiency rating based on points, points allowed, fouls, turnovers, etc. Players can be classified by their efficiency, allowing to balanced lineups throughout the game.

2 Question 6

A parametric approach relies on assumption about the population: - It follows a uniform distribution - The variance of the population is the same over each independent observation

An advantage is that the approach is simpler and generally will not overfit the training data. A disadvantage is that the model is sensitive to the assumptions made, making it inflexible and unfit to find relationships in nonlinear data.

A non-parametric approach doesn't make assumptions about the population. Random samples are used to make observations about the population.

An advantage is that the approach is not heavily affected by outliers, allowing for nonlinear data to be a good fit for these models. Variance decreases as the sample size increases and a drawback of this approach is the opposite; the variance of the model will be larger if the sample size is small.

3 Question 7

```
[2]: table = pd.read_csv('./q7Table.csv')
      table
```

```
[2]:
```

	X1	X2	X3	Y
0	0	3	0	Red
1	2	0	0	Red
2	0	1	3	Red
3	0	1	2	Green
4	-1	0	1	Green
5	1	1	1	Red

4 (a)

```
[3]: # if we were to compare to a non-origin point...
      # euclidean(p1, p2): array / object, array / object: {
      #   return math.sqrt((p1.x1-p2.x1)**2+(p1.x2-p2.x2)**2+(p1.x3-p2.x3)**2)
      # }
      def euclidean(x1, x2, x3):
          return math.sqrt(x1**2 + x2**2 + x3**2)

      arr = []
      for i in range(6):
          dist = euclidean(table['X1'][i], table['X2'][i], table['X3'][i])
          arr.append((i, dist))
      print(f'Obs: {i+1}, Dist: {dist}')
```

```
Obs: 1, Dist: 3.0
Obs: 2, Dist: 2.0
Obs: 3, Dist: 3.1622776601683795
Obs: 4, Dist: 2.23606797749979
Obs: 5, Dist: 1.4142135623730951
Obs: 6, Dist: 1.7320508075688772
```

```
[4]: prediction = []
      def knn(k):
          arr.sort(key=lambda x: x[1]) # Sort by minDist
          knn = arr[:k] # Add the k smallest distances to knn (obs: dist)

          # For each corresponding point, add to prediction and return the points
          for i in range(len(knn)):
```

```

        prediction.append(table.iloc[knn[i][0]])
    return prediction

# My knn function adds duplicates to prediction before returning?
# knn(3) is of len 4
# knn(4) is of len 9

```

5 (b)

```
[5]: knn(1)
```

```

[5]: [X1      -1
      X2       0
      X3       1
      Y      Green
      Name: 4, dtype: object]

```

The prediction is (-1, 0, 1), Green because this is the point closest to the origin

6 (c)

```
[6]: knn(3)
```

```

[6]: [X1      -1
      X2       0
      X3       1
      Y      Green
      Name: 4, dtype: object,
      X1      -1
      X2       0
      X3       1
      Y      Green
      Name: 4, dtype: object,
      X1       1
      X2       1
      X3       1
      Y      Red
      Name: 5, dtype: object,
      X1       2
      X2       0
      X3       0
      Y      Red
      Name: 1, dtype: object]

```

The prediction is red because the majority (2/3) of neighbors are colored red.

7 (d)

Given the boundary is highly non-linear, we should expect a small K to be the best fit. Our model is not flexible given the small observation size.

8 Question 9

```
[7]: Auto = pd.read_csv('./Auto.csv')
Auto.head()
```

```
[7]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70	
1	15.0	8	350.0	165	3693	11.5	70	
2	18.0	8	318.0	150	3436	11.0	70	
3	16.0	8	304.0	150	3433	12.0	70	
4	17.0	8	302.0	140	3449	10.5	70	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

```
[8]: Auto.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 397 entries, 0 to 396
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             397 non-null   float64
1   cylinders       397 non-null   int64
2   displacement    397 non-null   float64
3   horsepower      397 non-null   object
4   weight          397 non-null   int64
5   acceleration    397 non-null   float64
6   year            397 non-null   int64
7   origin          397 non-null   int64
8   name            397 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.0+ KB
```

9 (a)

Qualitative: year, origin, name, cylinders

Quantitative: mpg, displacement, hp, weight, acceleration

10 (b)

```
[9]: quantitative = []
quantitative.append(Auto[['mpg']])
quantitative.append(Auto[['displacement']])
quantitative.append(Auto[['weight']])
quantitative.append(Auto[['acceleration']])
```

```
[10]: def _range(arr, s):
        print(f'Range of {s}: [{arr.min()}, {arr.max()}] => {arr.max() - arr.
        ↪min()}')

        _range(Auto['mpg'], 'mpg')
        _range(Auto['displacement'], 'displacement')
        _range(Auto['weight'], 'weight')
        _range(Auto['acceleration'], 'acceleration')
```

```
Range of mpg: [9.0, 46.6] => 37.6
Range of displacement: [68.0, 455.0] => 387.0
Range of weight: [1613, 5140] => 3527
Range of acceleration: [8.0, 24.8] => 16.8
```

11 (c)

```
[11]: def _compute(arr, s):
        print(f'{s} => Std:{arr.std():.2f}\tMean:{arr.mean():.2f}')

        _compute(Auto['mpg'], 'mpg')
        _compute(Auto['displacement'], 'displacement')
        _compute(Auto['weight'], 'weight')
        _compute(Auto['acceleration'], 'acceleration')
```

```
mpg => Std:7.83 Mean:23.52
displacement => Std:104.38 Mean:193.53
weight => Std:847.90 Mean:2970.26
acceleration => Std:2.75 Mean:15.56
```

12 (d)

```
[12]: auto = Auto
auto.drop(index=[i for i in range(10,86,1)], inplace=True)
auto.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 321 entries, 0 to 396
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             321 non-null   float64
1   cylinders       321 non-null   int64
2   displacement    321 non-null   float64
3   horsepower      321 non-null   object
4   weight          321 non-null   int64
5   acceleration    321 non-null   float64
6   year           321 non-null   int64
7   origin          321 non-null   int64
8   name           321 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 25.1+ KB
```

```
[13]: def _range(arr, s):
        print(f'Range of {s}: [{arr.min()}, {arr.max()}] => {arr.max() - arr.
        ↪min()}')

        _range(auto['mpg'], 'mpg')
        _range(auto['displacement'], 'displacement')
        _range(auto['weight'], 'weight')
        _range(auto['acceleration'], 'acceleration')
```

```
Range of mpg: [11.0, 46.6] => 35.6
Range of displacement: [68.0, 455.0] => 387.0
Range of weight: [1649, 4997] => 3348
Range of acceleration: [8.5, 24.8] => 16.3
```

```
[14]: def _compute(arr, s):
        print(f'{s} => Std:{arr.std():.2f}\tMean:{arr.mean():.2f}')


        _compute(auto['mpg'], 'mpg')
        _compute(auto['displacement'], 'displacement')
        _compute(auto['weight'], 'weight')
        _compute(auto['acceleration'], 'acceleration')
```

```
mpg => Std:7.90 Mean:24.44
displacement => Std:99.86      Mean:187.17
weight => Std:809.64      Mean:2933.18
acceleration => Std:2.71      Mean:15.71
```

13 (e)

```
[15]: Auto.plot.scatter(x='weight',y='mpg')
```

```
[15]: <AxesSubplot:xlabel='weight', ylabel='mpg'>
```



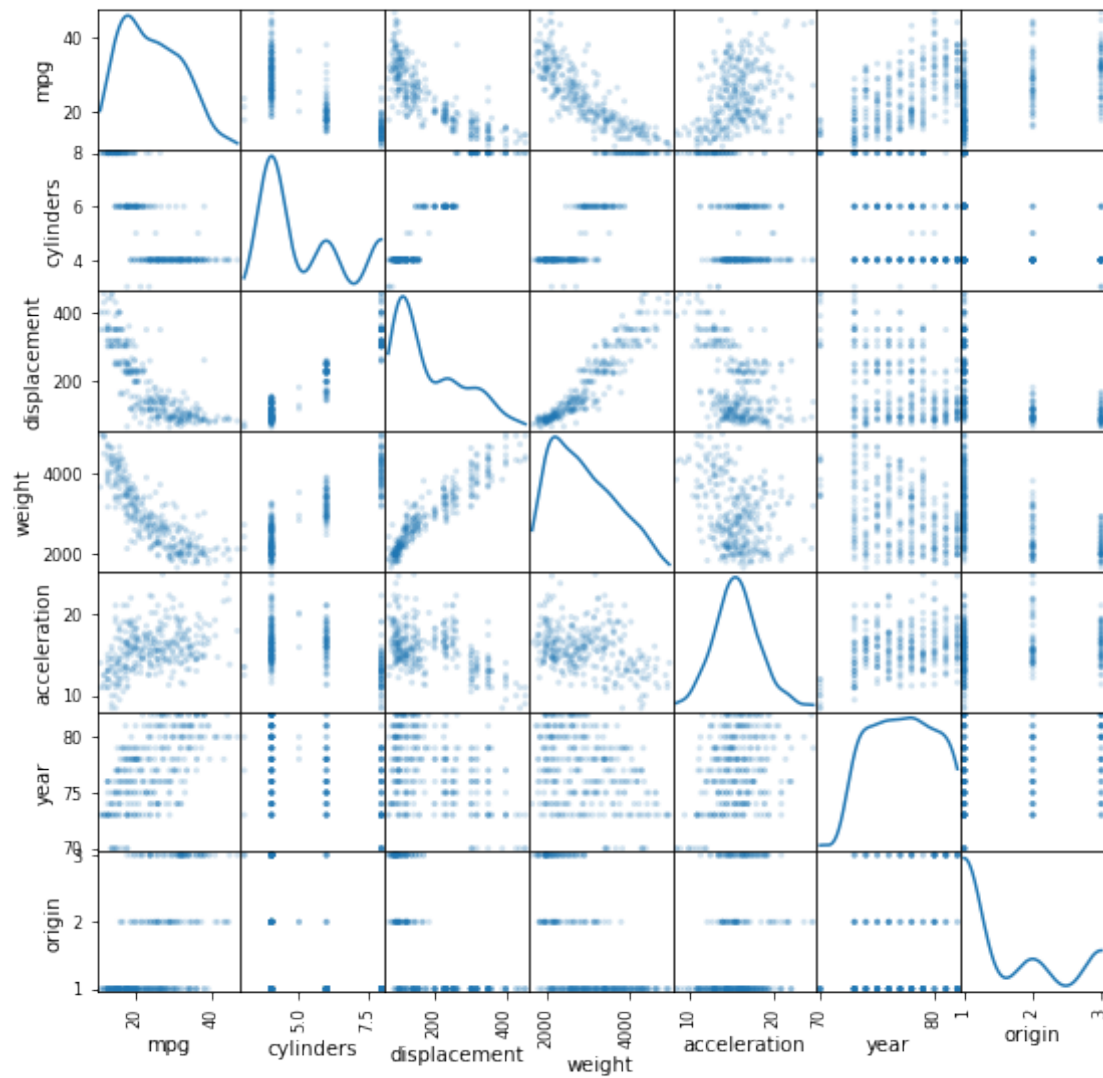
```
[16]: pd.plotting.scatter_matrix(Auto, alpha=0.2, figsize=(9,9), diagonal='kde')
```

```
[16]: array([[<AxesSubplot:xlabel='mpg', ylabel='mpg'>,
          <AxesSubplot:xlabel='cylinders', ylabel='mpg'>,
          <AxesSubplot:xlabel='displacement', ylabel='mpg'>,
          <AxesSubplot:xlabel='weight', ylabel='mpg'>,
          <AxesSubplot:xlabel='acceleration', ylabel='mpg'>,
          <AxesSubplot:xlabel='year', ylabel='mpg'>,
          <AxesSubplot:xlabel='origin', ylabel='mpg'>],
         [<AxesSubplot:xlabel='mpg', ylabel='cylinders'>,
          <AxesSubplot:xlabel='cylinders', ylabel='cylinders'>,
          <AxesSubplot:xlabel='displacement', ylabel='cylinders'>,
          <AxesSubplot:xlabel='weight', ylabel='cylinders'>,
          <AxesSubplot:xlabel='acceleration', ylabel='cylinders'>,
          <AxesSubplot:xlabel='year', ylabel='cylinders'>,
          <AxesSubplot:xlabel='origin', ylabel='cylinders'>],
         [<AxesSubplot:xlabel='mpg', ylabel='displacement'>,
          <AxesSubplot:xlabel='cylinders', ylabel='displacement'>,
          <AxesSubplot:xlabel='displacement', ylabel='displacement'>],
         ...])
```

```


<AxesSubplot:xlabel='weight', ylabel='displacement'>,
<AxesSubplot:xlabel='acceleration', ylabel='displacement'>,
<AxesSubplot:xlabel='year', ylabel='displacement'>,
<AxesSubplot:xlabel='origin', ylabel='displacement'>],
[<AxesSubplot:xlabel='mpg', ylabel='weight'>,
<AxesSubplot:xlabel='cylinders', ylabel='weight'>,
<AxesSubplot:xlabel='displacement', ylabel='weight'>,
<AxesSubplot:xlabel='weight', ylabel='weight'>,
<AxesSubplot:xlabel='acceleration', ylabel='weight'>,
<AxesSubplot:xlabel='year', ylabel='weight'>,
<AxesSubplot:xlabel='origin', ylabel='weight'>],
[<AxesSubplot:xlabel='mpg', ylabel='acceleration'>,
<AxesSubplot:xlabel='cylinders', ylabel='acceleration'>,
<AxesSubplot:xlabel='displacement', ylabel='acceleration'>,
<AxesSubplot:xlabel='weight', ylabel='acceleration'>,
<AxesSubplot:xlabel='acceleration', ylabel='acceleration'>,
<AxesSubplot:xlabel='year', ylabel='acceleration'>,
<AxesSubplot:xlabel='origin', ylabel='acceleration'>],
[<AxesSubplot:xlabel='mpg', ylabel='year'>,
<AxesSubplot:xlabel='cylinders', ylabel='year'>,
<AxesSubplot:xlabel='displacement', ylabel='year'>,
<AxesSubplot:xlabel='weight', ylabel='year'>,
<AxesSubplot:xlabel='acceleration', ylabel='year'>,
<AxesSubplot:xlabel='year', ylabel='year'>,
<AxesSubplot:xlabel='origin', ylabel='year'>],
[<AxesSubplot:xlabel='mpg', ylabel='origin'>,
<AxesSubplot:xlabel='cylinders', ylabel='origin'>,
<AxesSubplot:xlabel='displacement', ylabel='origin'>,
<AxesSubplot:xlabel='weight', ylabel='origin'>,
<AxesSubplot:xlabel='acceleration', ylabel='origin'>,
<AxesSubplot:xlabel='year', ylabel='origin'>,
<AxesSubplot:xlabel='origin', ylabel='origin'>]], dtype=object)

```

```
[17]: Auto['displacement'].plot.box()
```


```
[17]: <AxesSubplot:>
```



```
[18]: # Linear regression model: train weight v mpg
# Intercept: 48.306516:.06f Coefficient: -0.008315:.06f
from sklearn.linear_model import LinearRegression
import seaborn as sns
import matplotlib.pyplot as plt

X = Auto['weight'].values.reshape(-1,1)
y = Auto['mpg']
model = LinearRegression()
model.fit(X, y)

plt.figure(figsize=(6,6))
sns.regplot(data=Auto, x='weight', y='mpg')
plt.show()
```




```
[19]: ## Prediction: 4000lb car gets ~15.8mpg
model.predict(np.array([2000]).reshape(1,-1))

## Prediction: 2000lb car gets ~32mpg
model.predict(np.array([2000]).reshape(1,-1))
```

```
[19]: array([32.03637571])
```

```
[150]: sns.jointplot(data=Auto, x='weight', y='mpg', kind='reg', height=6)
plt.show()
```

```
[150]: <seaborn.axisgrid.JointGrid at 0x1aa0f7ad310>
```



[]: